

SELinux

DAC vs MAC:

Security-Enhanced Linux (SELinux, <https://github.com/SELinuxProject>) is a kernel module that implements a Mandatory Access Control (MAC) system. It was developed as a replacement for Discretionary Access Control (DAC) that ships with most Linux distributions with the form of the standard permissions/ACLs model.

If you are only using a DAC system :

*Administrators have no way to control users: a user could set world readable permissions on sensitive files (such as ssh keys). They could even `chmod +rwx` their entire home directory and nothing would stop them. The same reason can be extended to any (maybe trojaned) program executed by that user because processes inherit user's right. In summary: a regular user can grant (and restrict) access to their owned files to other users and groups or even change the owner of the file, leaving critical files exposed to accounts who don't need this access. It's true that this regular user could also restrict to be more secure, but that's discretionary: there's no way for the system administrator to enforce it for every single file in the system.

*Root (or sudo) access on a DAC system gives the (maybe hacked) person or the (maybe trojaned) program permission to perform as desired on a machine. In essence, there are two privilege levels (root and user), and no easy way to enforce a model of least-privilege.

But if you are using a MAC system, users won't be able to work around the rules previously set by the system administrator. These rules define what a user or process can do, confining every process to its own "domain" so the process can interact with only certain types of files and other processes from allowed domains. For instance, the Apache process can access only the `/var/www/html` directory and nothing more because there's a SELinux policy which mandates that.. This restriction is implemented from the kernel level and it's enforced as the SELinux policy loads into memory, so the access control becomes mandatory.

In SELinux everything is denied (even to root user): a series of exceptions policies must be written by sysadmin to give each element of the system (a service, process type or user) only the access required to function (to specific files, ports, pipes, ...). If a service, program or user subsequently tries to access or modify a file or resource not necessary for it to function, then this is denied and the action is logged. This prevents a hacker from hijacking any process to gain system-wide access

Because SELinux is implemented within the kernel, individual applications do not need to be especially written or modified to work under SELinux although, of course, if written to watch for the error codes which SELinux returns might work better afterwards. If SELinux blocks an action, this is reported to the underlying application as a normal (or, at least, conventional) "access denied" type error to the application.

Note SELinux policy is not something that replaces traditional DAC security. If a DAC rule prohibits a user access to a file, SELinux policy rules won't be evaluated because the first line of defense has already blocked access. SELinux security decisions come into play after DAC security has been evaluated. So if access to a resource is denied, please check access rights first. But note that if the DAC and MAC conflict, the SELinux policy takes priority. So, let's say that you change —as root— the ownership of httpd service to anyone (by running the command `chmod 777 httpd`); the default SELinux policy still prevents any random user from killing the web server.

A great introduction to SELinux is https://wiki.gentoo.org/wiki/SELinux/Quick_introduction and, in general, whole <https://wiki.gentoo.org/wiki/SELinux> site. To know even more, see: https://www.linuxtopia.org/online_books/getting_started_with_SELinux/SELinux_overview.html A very good introduction is https://access.redhat.com/documentation/en-us/red_hat_enterprise_linux/8/html/using_selinux/getting-started-with-selinux_using-selinux and https://www.server-world.info/en/note?os=CentOS_8&p=selinux&f=1

Activació de SELinux:

SELinux has two operating modes: permissive and enforcing. Permissive mode allows the system to function by using its regular DACs systems while logging every violation to SELinux. The enforcing mode enforces a strict denial of access to anything that isn't explicitly allowed by SELinux policies. You can see which mode system is running by executing this command:

getenforce

NOTA: Another command which can be used to know the same (though it shows other informational data) is *sestatus*. The meaning of the information shown by this command, among other interesting arguments, is explained in <https://www.thegeekstuff.com/2017/06/selinux-sestatus/>

You can immediately change to the desired mode simply executing...

sudo setenforce permissive (or *sudo setenforce enforcing*)

If you want this change to be permanent, you should modify the `"/etc/selinux/config"` file (which is linked by `"/etc/sysconfig/selinux"`). Specifically you should have this line...

SELINUX=permissive (or *SELINUX=enforcing* or now even *SELINUX=disabled*)

...and then rebooting the system.

If you enable SELinux after it has been disabled or you change the SELinux policy from the default targeted policy to another one (see NOTE below), it will be necessary to "relabel" the complete filesystem. To do this, execute this command...:

sudo touch /.autorelabel

...and then reboot the system (rebooting may take a while: that's because, in this circumstance, SELinux relabels every last file).

NOTA: By design, SELinux allows different policies to be written that are interchangeable by editing the *SELINUXTYPE=* line in `"/etc/sysconfig/selinux"` file. The default policy is the *"targeted"* policy which confines selected previously targeted system processes. By default many system processes are already targeted in RedHat family systems from its installation (including `httpd`, `named`, `dhcpd`, `mysqld`, etc) but all other system processes and all remaining userspace programs runs in an "unconfined domain" so they aren't covered by the SELinux protection model. One goal might be to target every process to force them to be run in a confined domain so the "targeted" policy protects as many key processes as possible, but it's responsibility from package maintainers. The alternative would be a deny-by-default model where every access is denied unless approved by the policy: it would be a very secure implementation, but this also means that developers have to anticipate every single possible permission every single process may need on every single possible object. So, in summary, the default behaviour sees SELinux concerned with only certain processes.

Conceptes bàsics de SELinux:

An SELinux policy defines which SELinux users belongs to a role, which roles are authorized to execute processes belonging to a domain and which domains are authorized to access to certain types of resources (a file, a link, a directory, a network port, a device, etc). In other words: a SELinux policy is a bunch of rules that say that so-and-so users can assume only so-and-so roles, and those roles will be authorized to execute only so-and-so domains' processes; the domains, in turn, can access only so-and-so file/port/... types. Note that:

*SELinux has a set of pre-built users. Every regular Linux user account is mapped to one or more SELinux users. For now we won't talk about them.

*A role is only used when SELinux is configured to implement RBAC (so it's not used by default). A role is a filter: it defines what users can assume the role in question and what domain this role can access to. For now we won't talk about them.

*A domain is the context within which a SELinux process can run. That context is like a wrapper ("a bubble") around the process: it tells the process what it can and can't do. In other words, the domain will define what files, directories, links, devices, or ports are accessible to the process, so it also can be thought as a "process type".

*A type is the context for a file that stipulates the file's purpose. For example, the context of a file may dictate that it's a web page, or that the file belongs to the "/etc" directory, etc.

Terminology tip 1: In general, when we call something as a SELinux "subject" we will be referring to a process that can potentially affect an object. An SELinux "object" is anything that can be acted upon: this can be a file, a directory, a port, a tcp socket, etc. The actions that a subject can perform on an object are the subject's permissions.

Terminology tip 2: The last step, where a process running within a particular domain can perform only certain operations on certain types of objects, is called Type Enforcement (TE).

Paràmetre -Z i comandes bàsiques SELinux: *chcon*, *semanage*, *restorecon*

As we have just said, the SELinux primary model or enforcement is called "type enforcement". Basically this means we define a "label" on a process to tie it to a domain, and a "label" on a file system object to tie it to a type, too. Policy rules then control access between labeled subjects (processes, daemons) and labeled objects (folders, files, devices, ports...). These labels (some times called "contexts" too) are stored as extended attributes on the file system or, for objects that aren't files (such as processes and ports) are managed by the kernel.

The label format is "user:role:type:sensitivity". Let's see the label/context of a file (for instance, the one of the "/etc/fstab" file) using the -Z parameter of ls command:

NOTA: The -Z switch will work with most utilities to show SELinux security contexts (*ps -eZ*, *id -Z*, *cp -Z*, *mkdir -Z*, *ss -Z*)

```
ls -Z /etc/fstab
```

We can see the SELinux security context fields: "system_u:object_r:etc_t:s0". As you can see, SELinux Users are suffixed by "_u", roles are suffixed by "_r" and types (for files) or domains (for processes) are suffixed by "_t".

*Each Linux user account maps to an SELinux user, and in this case, the root user that owns the file is mapped to the *system_u* SELinux user. This mapping is done by the SELinux policy.

*The second part specifies the SELinux role, which is *object_r*.

*What's most important here is the third part, the type of the file that's listed here as *etc_t*. This is the part that defines what type the file or directory belongs to. We can see that most files belong to the *etc_t* type in the "/etc" directory. We can also see some files may belong to other types, like *locale.conf* which has a *locale_t* type. Even when all the files listed here have the same user and group owners, their types could be different.

*The fourth part of the security context, *s0*, has to do with multilevel security or MLS. Basically this is another way of enforcing SELinux security policy different from "targeted" one, and this part shows the sensitivity of the resource (*s0*). We won't talk about this because it's a very advanced topic.

If we look at the SELinux security context of a file in our home directory we will see something like this: "unconfined_u:object_r:user_home_t:s0" (so we can infer that the type *user_home_t* is the default type for files in a user's home directory and a regular user is mapped by default to the *unconfined_u* SELinux user).

We can see the label of a (not running) binary, too: if we do `ls -Z /bin/ssh` we will see its context is "system_u:object_r:ssh_exec_t:s0". But if you want to know the context of a running process (for instance, the journald daemon), you could do...:

```
ps -eZ | grep "systemd-journald"
```

...and you will see the context is "system_u:system_r:syslogd_t:s0" (so its domain is "syslogd_t").

As we already know, thanks to a defined policy, each process domain can act on only certain types of files and nothing more so, even if a process is hijacked by another malicious process or user, the worst it can do is to damage the files it has access to. For instance, a process running in the `httpd_t` domain, for instance, can interact with an object labeled with the `httpd_something_t` type (because a defined policy says so). So, as Apache runs in the `httpd_t` domain it can read `/var/www/html/index.html` because it is of type `httpd_sys_content_t` but it can not access `/home/username/myfile.txt` even though this file is world readable because its context is not accessible by `httpd_t` type. If Apache were to be exploited, it would not be able to start any process not in the `httpd_t` domain (which prevents escalation of privileges) or access any file not in an `httpd_t` related domain.

Unless specified by the policy, processes and files are created with the contexts of their parents. So if we have a process called "proc_a" spawning another process called "proc_b", the spawned process will run in the same domain as "proc_a" unless specified otherwise by the SELinux policy. Similarly, if we have a directory with a type of "some_context_t", any file or directory created under it will have the same type context unless the policy says otherwise. However, this inheritance is not preserved when files are copied to another location: in a copy operation, the copied file or directory will assume the type context of the target location (though this change of context can be overridden by the `--preserve=context` clause in `cp` command).

We can use `chcon` command to change the context of a file/directory (in a similar way to how `chown` or `chmod` may be used to change the ownership or standard file permissions of a file). For instance, to change the context of all files inside `/home/dan/html` folder (-R parameter is for "recursively"), run this:

```
sudo chcon -vR -t httpd_sys_content_t /home/dan/html
```

NOTA: If you don't know the label but you know a file with the equivalent labeling you want, you can do:
`sudo chcon -v --reference=/path/file1 /path/file2`

NOTA: You can change all the context simply doing `chcon -v user:role:type:sensitivity /path/file` or any other of the parts of the context with `-u`, `-r` or `-l` arguments, respectively.

Modifying contexts via `chcon` will persist between system reboots but they will last only until that modified portion of the filesystem is relabeled. If you want a context change will be permanent, even through a complete filesystem is relabeled, you should use `semanage fcontext -a` and `restorecon` commands: the former writes the desired context inside the `/etc/selinux/targeted/contexts/files/file_contexts.local` file so it can "remembered" (and effectively applied) later by the latter. So this alternate and more "enduring" way to change the context of a file/directory is by first running...:

```
sudo semanage fcontext -a -t httpd_sys_content_t '/home/dan/html(/.*)?'
```

NOTA: Note the use of regular expressions to tell the content of the specified folder instead of a non-existing -R parameter

NOTA: Contexts of files already existing in the system are listed in the `/etc/selinux/targeted/contexts/files/file_contexts` file: it's a large file and it lists every file type associated with every application supported by the Linux distribution. Contexts of new directories and files modified by `semanage fcontext` command are recorded in the `/etc/selinux/targeted/contexts/files/file_contexts.local` file. So when we run the `restorecon` command, SELinux will look up the correct context from one of these two files and apply it to the target

NOTA: Instead of `-a` parameter, in `semanage fcontext` we could use `-m` (to modify) or `-d` (to delete) parameters

NOTA: If you don't know the label but you know a file with the equivalent labeling you want, you can do: `sudo semanage fcontext -a -e /path/file1 /path/file2` On the other hand, to check the context of all files/directories in the filesystem we can execute: `sudo semanage fcontext -l`

...and then, to recursively restore the default contexts stored in "file_contexts.local" file for all files inside the whole directory:

```
sudo restorecon -Rv /home/dan/html
```

NOTA: To restore just the index.html file, we would use `sudo restorecon -v /home/dan/html/index.html`

Another example where `semanage fcontext` and `restorecon` commands can help is the case of an user which edits a copy of an "index.html" file in his/her home directory (so it will have the `user_home_t` context) and moves (`mv`) the file to the `/var/www/html` folder (which has a different context by default: `httpd_sys_content_t`). Whilst the copy (`cp`) command will typically adopt, as we already know, the default destination directory's or file's security context, move (`mv`) maintains the source's security context. So Apache would fail to load this file because it won't have the correct label. We could use the `chcon` command to change the context of the file in question but we alternatively could just assign the (default) context specified in `/etc/selinux/targeted/contexts/files/file_contexts` or `file_contexts.local` files (which it will be `httpd_sys_content_t` because file in question is located inside `/var/www/html` and this folder -and all its content- is already specified there with this context). Besides, running `chcon` requires you to know the correct context for the file but `restorecon` doesn't need this specified.

NOTA: If we simply wanted to examine the contexts of the `/var/www/html` directory to see if any files needed their contexts restored, we can use `restorecon` with the `-n` switch to prevent any relabelling occurring: `restorecon -Rv -n /var/www/html`. Other arguments of `restorecon` are `-f listfilestochange.txt`, `-o listchangedfiles.txt` or `-e /subfolder/to/exclude`

NOTA: There is a nifty tool called `matchpathcon` that can help troubleshoot context-related problems. This command will look at the current context of a resource and compare it with what's listed under the "context database" (that is, the "file_contexts" files). If different, it will suggest the change required. Let's test this with the `/www/html/index.html` file: `matchpathcon -V /www/html/index.html` (the `-V` flag means "verify")

NOTA: Another interesting tool is `seinfo -t` (available from "setools-console" package) which lists all contexts currently in use on your system; `grep` for the name of your application.

If you want to allow some labeled processes (say, the Apache web server) to run in such a way that it only generates SELinux error messages rather than just failing, instead of using `setenforce permissive` (which puts all domains on the system into permissive mode) you can use following command instead, which puts just the domain specified):

```
sudo semanage permissive -a httpd_t
```

NOTA: To know what domain you should put into permissive mode (that is, the domain which has denied problems), you could look at the type in the context of the "scontext" field from audit log entries (that is, the ones with the field "type" with the "AVC" and "denied" values associated; more about this later)

Simultaneously, your other processes run under SELinux's restricted mode protection. Once you have that process running without any error messages, you can return it to enforcing mode with:

```
sudo semanage permissive -d httpd_t
```

NOTA: To check on the current SELinux mode status of your processes, run: `sudo semanage permissive -l`

Un altre "object" diferent de fitxers/directoris: ports

Network ports are labeled too. This means that, due to predefined policies, only processes belonging to specific domains will be able to listen on specific labeled ports. Or said backwards, only specific labeled ports will be able to be attached to specific domains. This means, for instance, that if we want to change the port where some server is listening (for instance, a SSH server by editing its `/etc/ssh/sshd_config` configuration file), we should label the new port with the same context old 22/TCP port had (in this case, the `ssh_port_t` context) in order to allow SSH server to "access" to this new port.

NOTA: A full list of port contexts bound to allowed port numbers can be obtained executing: `sudo semanage port -l`. Additionally we can use the `-C` flag, which only displays customizations (the full contents of `-l` are 400 lines!)

To label a unlabeled port (for instance, 2222/TCP) to a port type (in this case, `ssh_port_t`) we can do:

```
sudo semanage port -a -t ssh_port_t -p tcp 222
```

To undo this (or, in general, to delete a label from a port) we can do:

```
sudo semanage port -d -t ssh_port_t -p tcp 2222
```

NOTA: If we were instead meant to apply port 2222/TCP to Apache server, we can use the `-m` flag to modify the type from `ssh_port_t` to `http_port_t`: `sudo semanage port -m -t http_port_t -p tcp 2222`

NOTA: The `sepolicy` command, which comes with the `policycoreutils-devel` package, can also be used to view SELinux port types that have been set on a particular port.

Booleans SELinux

Booleans allows us to make changes to part of policy at runtime without need for having knowledge of policy writing. This allows changes to be implemented without the need to recompiling a SELinux policy or adding your own custom SELinux modules with `audit2allow` (see below). For example, let's say we want to share our user's home directory over FTP for read-write access and we have already shared them but while trying to access them, we can't see them: that's because SELinux policy is preventing the FTP daemon from reading and writing in user's home directory. We need change the policy so that ftp can access home directories, to do that we will see if there are any Booleans available to accomplish it by running:

```
sudo semanage boolean -l
```

It will produce a list of all available Booleans with their current status (on or off) and a description.

NOTE: You can also get list of available Booleans by running `getsebool -a` but it will not show its description

To turn on/off a Boolean you can use the `setsebool` command. For instance, if you execute this...:

```
sudo setsebool ftp_home_dir on
```

...now our ftp daemon will be able to access user's home directory.

Another example: the default behavior is that a web application running in the `httpd_t` context will not be allowed to send emails. That helps greatly to prevent a vulnerable web application to send out SPAM. Well, if you want to operate a web mail service Apache must be able to send emails. No big deal:

```
setsebool -P httpd_can_sendmail on
```

NOTA: With the `-P` option all pending values are written to the policy file on disk so they'll be persistent across reboots.

NOTA: Hi ha més trucs aquí: <https://wiki.centos.org/TipsAndTricks/SelinuxBooleans>. On the other hand, a full example interesting to read is <https://www.rootusers.com/deploy-a-basic-cgi-application-with-apache>

NOTA: Even when a system is in enforcing mode a malicious user with root access can manipulate policies or put SELinux into permissive mode. There is a method to prevent this: lock down your system doing `setsebool -P secure_mode_policyload on`. Be aware: Once active nothing can not be changed during runtime, you need to reboot your system and provide `selinux=1 enforcing=0` as grub boot parameters to be able to change any SELinux settings.

Troubleshooting SELinux

Sooner or later you may run into situations where SELinux denies access to something and you need to troubleshoot the issue. There are a number of fundamental reasons why SELinux may deny access to a file, process or resource:

- *A mislabeled file.
- *A process running under the wrong domain
- *A wrong policy tuning due to working with booleans.
- *A bug in policy (an application requires access to a file that wasn't anticipated when the policy was written and generates an error)
- *An intrusion attempt.

The first 4 we can deal with, whereas giving alarm and notice in the 5th case is exactly the intended behaviour. To troubleshoot any issue, the log files are key and SELinux is no different. By default SELinux log messages are written to `/var/log/audit/audit.log` via the *auditd* daemon, which is started by default. SELinux log messages are labeled with the "AVC" keyword so that they might be easily filtered from other messages with *grep* or with specific client commands provided by the *audit* subsystem itself, as we can see in next table:

Audit commands cheatsheet

Some of the most common searches to find AVC error messages are:

- *All error messages: ***ausearch -m avc***
- *Today's error messages: ***ausearch -m avc -ts today***
- *Error messages from the past 10 minutes: ***ausearch -m avc -ts recent***
- *Or, to find SELinux denials for a particular service, use the `-c comm` option, where "comm" is the executable's name
For example, httpd for the Apache web server or smbd for Samba: ***ausearch -m avc -c httpd***

An AVC error message can these fields (among others):

type=AVC : SELinux caches access control decisions for resource and processes. This cache is known as the Access Vector Cache (AVC). This field is saying the entry is coming from an AVC log and it's an AVC event.
denied { getattr }: The permission that was attempted and the result it got. In this case the "get attribute" operation was denied.
pid: The process id of the process that attempted the access.
comm: The process id by itself doesn't mean much. The comm attribute shows the process command (httpd, for instance)
path: The location of the resource that was accessed (for instance, a file under `/www/html/index.html`)
dev and **ino**: The device where the target resource resides and its inode address.
scontext: The security context of the source process's domain (for instance, *httpd_t* domain).
tcontext: The security context of the target resource (for instance, *user_home_t* type)
tclass: The class of the target resource (for instance, a file)

Per saber més sobre els possibles camps disponibles, consultar: <https://wiki.centos.org/HowTos/SELinux#head-c84fad68bffc5eca190fa3a1aab3cac2cfe94a63>

Besides the generic *ausearch* command, we can install the "setroubleshoot-server" package (and its GUI too, the "setroubleshoot" package) to enjoy the *sealert* command, which gives more information to debug the eventual AVC denials letting users to know when access has been denied, helping them to resolve it if necessary, and reducing overall frustration.

NOTA: The "setroubleshoot-server" package provides a service called *setroubleshootd* (activated by D-Bus, so "dbus" service should be running on system to that *setroubleshootd* works fine), of wich the *sealert* command acts as a "frontend".

NOTA: You can even notify yourself via email when some AVC happens with *setroubleshootd* server!. To do so, open `/etc/setroubleshoot/setroubleshoot.conf` file and adjust the [email] section to fit your server:

```
recipients_filepath = /var/lib/setroubleshoot/email_alert_recipients
smtp_port = 25
smtp_host = localhost
from_address = selinux@myserver.com
subject = [MyServer] SELinux AVC Alert
```

and then, execute this: `echo "destiny@mycompany.com" >> /var/lib/setroubleshoot/email_alert_recipients`

To see all available information about all "AVC denied" log messages saved in `/var/log/audit/audit.log`, you can execute...:

```
sudo sealert -a /var/log/audit/audit.log
```

...and you'll see a friendlier explanation of what went wrong than just having executed `ausearch` command. You'll obtain a suggestion for a fix, too.

NOTA: You can also search the message_ID with the AVC message number you are interested in and execute this to obtain even more detailed information: `sudo sealert -l [message_ID]`

In more complex cases maybe you'll need to use the `audit2why` command, which translates AVC messages into a description of why the access was denied (it's equivalent to execute `audit2allow -w`). So you could running this command to show at screen "what's wrong is happening":

```
sudo cat /var/log/audit/audit.log | audit2why (or sudo audit2why -a or sudo audit2allow -w -a)
```

Exploring default SELinux Policies

The SELinux policies for Fedora are shipped in the "selinux-policy" package. The policy rules come in the form: `"allow domain type:class { permission1 permission2 ... }; [boolean]:True"`, where:

- *Domain : Label for the process/es
- *Type : Label for the object/s
- *Class : The kind of object being accessed (e.g. "file", "socket")
- *Permission : The operation being performed (e.g. "read", "write")
- *Boolean : A boolean which would allow this rule if it was "on"

We can use the `sesearch` command (available from "setools-console" package) to check, for instance, the type of access allowed for the httpd daemon to files of `httpd_sys_content_t` type:

```
sesearch --allow -s httpd_t -t httpd_sys_content_t -c file
```

Notice the first line: `"allow httpd_t httpd_sys_content_t : file { ioctl read getattr lock open };"` It says that the httpd daemon has I/O control, read, get attribute, lock, and open access to files of the `httpd_sys_content_t` type. We could "filter" even more adding the `-p` argument to specify a permission of which we are interested (it can specified multiple times).

Another real example of a httpd-related rule is: `"allow httpd_t httpd_log_t: file { append create getattr ioctl lock open read setattr };"` The rule allows any process labeled as `httpd_t` to create, append, read and lock files labeled as `httpd_log_t`. Remember using the `ps` command, you can list all processes with their labels (`ps -eZ | grep httpd`) and using `semanage` you can know which objects are labeled as `httpd_log_t` (`semanage fcontext -l | grep httpd_log_t`).

In general, we could collate `sesearch` output with the AVC-denial records in the `auditd` log to infer the necessary custom policy that would be necessary to add (see next chapter) in order to allow a requested access to the target.

`Sesearch` can be useful too to identify any rule that is toggled by booleans: the identifier shown in square brackets is the name of the boolean that would allow this access if it was "true". On the other hand, if we wanted to inspect exactly what any boolean is allowing (that is, which rules are equivalent to activate any boolean), we could do this:

```
sesearch --allow -b boolean_name [-c class] [-p permission]
```


Creating a custom SELinux Policy

If simply changing a context via *chcon/semanage/restorecon* is not sufficient because the denial case is too complex, you can try to add custom rules to your policy. The way to do it is by using the *audit2allow* command, which serves to generate "allow/dontaudit" policy rules from logs of denied operations. To implement it (first on screen) you can simply *grep* for the misbehaving process (for instance, a "httpd" server) and pass these resulting AVC logs to *audit2allow* like this:

```
sudo grep httpd_t /var/log/audit/audit.log | sudo audit2allow -m my_policy_name  
(or sudo ausearch -m avc -c httpd --raw | sudo audit2allow -m my_policy_name)
```

Review the result to check if it makes sense (ensure your *grep* statement does not catch too much). If you're confident it's okay, run the same command again with a capital *-M* as parameter: it instructs the command to create a type enforcement (.te) file with the specified name (there you will be able to view the rules to be added) but it also compiles the rule into a SELinux policy package (.pp) ...

```
sudo ausearch -m avc -c httpd --raw | sudo audit2allow -M my_policy_name
```

NOTA: Do not use the "-M" option to specify the same module name more than once. For example, if you run the command below once with "-M local", and want to run it again later, choose a different name, such as "-M local2".

...which then can be installed with the *semodule* command:

```
sudo semodule -i my_policy_name.pp
```

When an SELinux-enabled system starts, the policy (that is, the set of rules) is loaded into memory. SELinux policy comes in modular format so rules are grouped by modules and modules can be dynamically added/removed/reloaded/upgraded/enabled/disabled from memory at run time (using the *semodule* command, by the way). The SELinux "policy store" (whose name is shown by *sestatus* command) keeps track of the modules that have been loaded and we can use the *semodule -l* command to list the SELinux policy modules currently loaded into memory:

```
sudo semodule -l | less
```

Most modern distributions include binary versions of the modules as part of the SELinux packages; they are located inside *"/var/lib/selinux/targeted/active/modules"* folder : if you look closely, they will relate to different applications (but they're not human-readable, though).

The way SELinux modularization works is that when the system boots, policy modules are combined into what's known as the active policy. This policy is then loaded into memory. The combined binary version of this loaded policy can be found under the *"/etc/selinux/targeted/policy"*

EXERCICIS:

1.-Arrenca una màquina virtual Fedora Workstation (ha de ser Fedora, no val Ubuntu) i fes les següents accions:

NOTA: Ha de ser Fedora perquè les úniques distribucions que implementen SELinux són les de la família de RedHat. Les distribucions Ubuntu/Debian/Suse fan servir un altre tipus de mòdul LSM (Linux Security Module) anomenat AppArmor (<https://wiki.ubuntu.com/AppArmor>), no tan flexible ni detallat com SELinux.

a) Instal·la el paquet "httpd" (que representa el servidor web Apache) i digues quin és el contexte SELinux de la carpeta "/var/www/html" (i quina comanda has fet servir per saber-ho).

b) Inicia el servei "httpd" (`sudo systemctl start httpd`) i digues en quin domini SELinux s'executa per defecte (i quina comanda has fet servir per saber-ho).

c) Instal·la el paquet "setools-console" per tal de poder executar la comanda `sesearch --allow -s httpd_t -t XXX -c file` (on les "X" han de substituir-se pel tipus SELinux mostrat a la sortida de la comanda `ls -Z /home`) i així deduir, amb la sortida que observis, si el servidor Apache serà capaç d'accedir/llegir (i per tant, de publicar) algun fitxer que estigui guardat a la teva carpeta personal. ¿Sí o no? Al següent exercici ho comprovaràs

Per a què un servidor Apache pugui accedir/llegir el contingut d'una determinada carpeta s'han de complir tres requisits:

*La seva configuració interna ho ha de poder permetre. Això es pot fer afegint dins del seu arxiu de configuració "/etc/httpd/conf.d/httpd.conf" una secció tal com aquesta:

```
<Directory "/ruta/carpeta">
Require all granted
</Directory>
```

*Els permisos de la carpeta han de ser "r" i "x" per l'usuari/grup amb el què s'executa el servidor Apache, el qual s'anomena "apache" en el cas de sistemes RedHat/Fedora i "www-data" en el cas de sistemes Debian/Ubuntu

*El domini del procés `httpd` ha de tenir assignat una regla que li permeti accedir a carpetes/fitxers etiquetats amb el context adient

2.-a) Modifica l'arxiu "/etc/httpd/conf.d/httpd.conf" per tal de substituir la línia `DocumentRoot "/var/www/html"` per aquesta altra: `DocumentRoot "/home/usuari"` (o la ruta de la carpeta personal de l'usuari que estiguis fent servir) i afegir-ne a continuació la següent secció:

```
<Directory "/home/usuari">
Require all granted
</Directory>
```

b) Executa la comanda `chmod -R 755 /home/usuari` per tal de què l'usuari "apache" (i, de fet, qualsevol altre) pugui accedir i llegir el contingut de la teva carpeta personal.

c) En un sistema que no tinguis SELinux en mode "Enforcing", ja n'hi hauria prou per a què l'Apache pogués oferir contingut ubicat dins de "/home/usuari". Comprovem que això en el nostre cas no és així executant la comanda `echo "Hola" > /home/usuari/pagina.txt`, reiniciant el servidor `httpd` i obrint un navegador (a la mateixa màquina virtual mateixa) per anar a la direcció <http://127.0.0.1/pagina.txt> ¿Què veus? Per què?

d) Ara canvia el context de l'arxiu `pagina.txt` fent servir la comanda `chcon` per tal de què sigui el mateix que el del contingut de la carpeta "/var/www/html". Comprova que, efectivament, ho hakis aconseguit fent `ls -Z pagina.txt` i torna tot seguit a anar a la direcció <http://127.0.0.1/pagina.txt> (no cal reiniciar el servidor per res) ¿Què veus? Per què?

3.-a) Executa la comanda `sudo ausearch -m AVC` Hauràs de veure el missatge d'error generat per SELinux provocat per l'apartat c) de l'exercici anterior. ¿On es troba físicament guardat aquest missatge? ¿Què significa el valor "read" escrit entre claus? ¿I els valors "pid", "comm", "name", "scontext", "tcontext", "tclass" i "permissive"?

b) Instal·la el paquet "setroubleshoot-server" per tal de poder executar la comanda `sudo sealert -a /var/log/audit/audit.log`. Veuràs que et diagnostica el mateix problema vist a l'apartat anterior però a més proposa diverses solucions. ¿Què fa, en concret, la comanda suggerida `sudo setsebool -P httpd_read_user_content 1`? Executa-la i observa a continuació la sortida de la comanda `getsebool httpd_read_user_content` per comprovar que ho hagi fet bé. Tot seguit...

c) ...executa ara la comanda `echo "Adeu" > /home/usuari/pagina2.txt` i obre de nou el navegador per tal d'anar a la direcció <http://127.0.0.1/pagina2.txt>. ¿Què veus? ¿Per què penses que no ha calgut ara canviar l'etiqueta de "pagina2.txt" encara que aquest continui sent "user_home_t" per tal de què la pàgina es vegi correctament? Per confirmar la resposta que acabes de donar a aquesta pregunta, observa la sortida de la comanda `sesearch --allow -s httpd_t -c file | grep httpd_user_read_user_content` i dedueix si el que et mostra té coherència o no amb aquesta resposta.

d) Una altra solució suggerida per la comanda `sealert` feta servir a l'apartat b) era executar aquestes comandes: `sudo ausearch -c "httpd" --raw | audit2allow -M my-httpd && sudo semodule -X 300 -i my-httpd.pp`. ¿Què és l'arxiu "my-httpd.pp" que genera la comanda `audit2allow`? ¿Per a què serveix la comanda `semodule`?

e) ¿Què passa si ara copies l'arxiu "/etc/fstab" dins de la teva carpeta personal executant exactament aquesta comanda: `cp --preserve=context /etc/fstab /home/usuari` i proves seguidament d'obrir al navegador la direcció <http://127.0.0.1/fstab>? ¿Per què? ¿Què passa en canvi si tornes a accedir a la mateixa direcció però havent copiat abans el mateix fitxer a la mateixa carpeta però executant ara la comanda `cp /etc/fstab /home/usuari`? Confirma la resposta a ambdós casos amb l'ajuda de la comanda `ls -Z /home/usuari/fstab`

4.-a) Ara canviarem el port per on escoltarà el servidor Apache. Per aconseguir això hem de modificar la línia `Listen 80` de l'arxiu "/etc/httpd/conf/httpd.conf" per `Listen 12345` i reiniciar el servei. No obstant, en voler fer aquest darrer pas veurem que passa un error (degut a què tenim SELinux en mode "enforcing"). Torna a executar la comanda `sudo sealert -a /var/log/audit/audit.log` per confirmar que, efectivament, és un error AVC de tipus "name_bind" (és a dir, de vinculació a un port).

b) Executa la comanda `semanage port` adient per tal d'aconseguir que el port 12345 estigui etiquetat amb el context adient (l'hauràs d'esbrinar, per exemple amb la comanda `semanage port -l`) per tal de què el domini `httpd_t` pugui accedir-hi. Un cop fet això, comprova que ja podràs iniciar el servei correctament i que aquest estarà escoltant al port 12345 sense problemes (això últim ho pots comprovar, per exemple, observant la sortida de la comanda `ss -tnl`)

NOTA: Gràcies a tenir el servei setroubleshoot-server ja funcionant, el missatge d'error s'envia no només a l'arxiu "audit.log" sinó també al registre d'events general del sistema (accessible via la comanda genèrica `journalctl`, com veurem en temes posteriors).

SELinux (II)

SELinux users and roles:

Roles are not that important when we see them in file security contexts. For files, it's listed with a generic value of `object_r`. Roles become important when dealing with users and processes. SELinux users are different entities from normal Linux user accounts, including the root account. An SELinux user is not something you create with a special command, nor does it have its own login access to the server. Instead, SELinux users are defined in the policy that's loaded into memory at boot time, and there are only a few of these users. The user names end with `_u`, just like types or domain names end with `_t` and roles end with `_r`. Different SELinux users have different rights in the system and that's what makes them useful.

The SELinux user listed in the first part of a file's security context is the user that owns that file. This is just like you would see a file's owner from a regular `ls -l` command output. A user label in a process context shows the SELinux user's privilege the process is running with.

When SELinux is enforced, each regular Linux user account is mapped to an SELinux user account. There can be multiple user accounts mapped to the same SELinux user. This mapping enables a regular account to inherit the permission of its SELinux counterpart. To view this mapping, we can run following command:

```
sudo semanage login -l
```

The first column in this table, "Login Name", represents the local Linux user accounts. Any regular Linux user account is first mapped to the "default" login. This is then mapped to the SELinux user called `unconfined_u` (this is the second column). The third column shows the multilevel security / Multi Category Security (MLS / MCS) class for the user but for now, let's ignore that part and also the column after that (Service). Next, we have the root user. Note that it's not mapped to the "default" login, rather it has been given its own entry. Once again, root is also mapped to the `unconfined_u` SELinux user.

'system_u' is a different class of user, meant for running processes or daemons. To see what SELinux users are available in the system, we can run following command:

```
sudo semanage user -l
```

What does this bigger table mean? First of all, it shows the different SELinux users defined. We had seen users like `unconfined_u` and `system_u` before, but we are now seeing other types of users like `guest_u`, `staff_u`, `sysadm_u`, `user_u` and so on (the names are somewhat indicative of the rights associated with them). Let's look at the fifth column, SELinux Roles; if you remember from the first part of this tutorial, SELinux roles are like gateways between a user and a process. We also compared them to filters: a user may enter a role, provided the role grants it. If a role is authorized to access a process domain, the users associated with that role will be able to enter that process domain.

Now from this table we can see the `unconfined_u` user is mapped to the `system_r` and `unconfined_r` roles. Although not evident here, SELinux policy actually allows these roles to run processes in the `unconfined_t` domain. Similarly, user `sysadm_u` is authorized for the `sysadm_r` role, but `guest_u` is mapped to `guest_r` role. Each of these roles will have different domains authorized for them.

Now if we take a step back, we also saw from the first code snippet that the default login maps to the `unconfined_u` user, just like the root user maps to the `unconfined_u` user. Since the `**default**` login represents any regular Linux user account, those accounts will be authorized for `system_r` and `unconfined_r` roles as well. So what this really means is that any Linux user that maps to the `unconfined_u` user will have the privileges to run any app that runs within the `unconfined_t` domain (as you can imagine, unconfined processes would have all types of access in the system but even then, this full access is not arbitrary: full access is also specified in the SELinux policy). To demonstrate this, let's run this command as a regular user:

```
id -Z
```

This shows the SELinux security context for root: "unconfined_u:unconfined_r:unconfined_t:s0-s0:c0.c1023". So the root account is mapped to the *unconfined_u* SELinux user, and *unconfined_u* is authorized for the *unconfined_r* role, which in turn is authorized to run processes in the *unconfined_t* domain. If you execute the same *id -Z* command for the root user, you will get the same output: a regular user maps to the same SELinux user/role/domain than the root account also maps to.

We had seen the list of a number of SELinux users before:

**guest_u*: This user doesn't have access to X-Window system (GUI) or networking and can't execute *su* / *sudo* command.

**xguest_u*: This user has access to GUI tools and networking is available via Firefox browser.

**user_u*: This user has more access than the guest accounts (GUI and networking), but can't switch users by running *su* or *sudo*.

**staff_u*: Same rights as *user_u*, except it can execute *sudo* command to have root privileges.

**system_u*: This user is meant for running system services and not to be mapped to regular user accounts.

To see how SELinux can enforce security for user accounts, let's think about a "regularuser" account. As a system administrator, you now know the user has the same unrestricted SELinux privileges as the root account and you would like to change that. Specifically, you don't want the user to be able to switch to other accounts, including the root account. So we will change regularuser's SELinux user mapping to *user_u*:

```
sudo semanage login -a -s user_u regularuser
```

The change won't take effect until regularuser logs out and logs back in. If we now run the *id -Z* command again (or *semanage login -l*) to see the new SELinux context for regularuser.

NOTA: We can assign a SELinux user to a group user simply writing the group with a first "%" character, like this: *sudo semanage login -a -s user_u %regularusers*

Let's see another example of restricting user access through SELinux. By default, SELinux allows users mapped to the *guest_t* account to execute scripts from their home directories. We can run the *getsebool* command to check the boolean value (by default is on) like this: *getsebool allow_guest_exec_content*. If we change the SELinux user mapping for the "regularuser" account executing *sudo semanage login -a -s guest_u regularuser* and turn off former boolean executing *sudo setsebool allow_guest_exec_content off*, we will see "regularuser" can't execute any script from their home folder. So this is how SELinux can apply an additional layer of security on top of DAC. Even when the user has full read, write, execute access to the script created in their own home directory, they can still be stopped from executing it

Let's see another example of restricting user access through SELinux. What if you want to stop this particular user from starting the *httpd* service even when the user's account is listed in the *sudoers* file? First, execute *sudo semanage login -a -s user_u regularuser*. Now that regularuser has been restricted to *user_u* (and that means to role *user_r* and domain *user_t*), we can verify its access executing *sudo semanage user -l* or this other new command:

```
sudo seinfo -u user_u -x
```

The output shows the roles *user_u* can assume (it is *user_r*). Taking it one step further, we can run the *seinfo* command to check what domains the *user_r* role is authorized to enter:

```
sudo seinfo -r user_r -x
```

There are a lot of domains *user_r* is authorized to enter but ...does this list show *httpd_t* as one of the domains? There are a number of *httpd* related domains the role has access to, but *httpd_t* is not one of them. So, if the regularuser account tries to start the *httpd* daemon, the access should be denied because the *httpd* process runs within the *httpd_t* domain and that's not one of the domains the *user_r* role is authorized to

access. And we know *user_u* (mapped to regularuser) can assume *user_r* role. This should fail even if the regularuser account has been granted sudo privilege.

SELinux sandbox

As the name implies sandbox is used to run any command in a fully user-controllable SELinux-enforced sandbox. This allows you to isolate commands/applications from the rest of your system and only grant specific permissions and capabilities. Even better, SELinux includes a handful of prebuilt sandbox types that allow access to certain critical resources. After you install one dependency (policycoreutils-sandbox), you can even run X apps in an SELinux sandbox! Here's the magic command for running Firefox in an SELinux sandbox on Fedora:

```
sandbox -X -t sandbox_net_t -t sandbox_web_t -w 1280x1024 firefox
```

This runs a sandbox with its own X server (-X), allows ports required for web browsing and general network access (-t sandbox_web_t and -t sandbox_net_t) and launches firefox in a 1280x1024 window (-w 1280x1024). This will open up a new window with a completely clean instance of Firefox that is isolated from the rest of your processes by SELinux. Note that this also means you won't be able to access any of your files (including your Firefox profile) so you will get a completely fresh session every time.

If you find this handy, you might want to try extending this to other apps that you use where you might want to test things in clean environments or are handling files you don't entirely trust. For example, here's how to open a report.pdf from my home directory in a sandboxed PDF viewer:

```
sandbox -X -w 1280x1024 -i ~/report.pdf evince report.pdf
```

Multilevel security (MLS)

MLS is the fine-grained part of an SELinux security context. It gives the concept of hierarchical "sensitivity". By hierarchy, we mean the levels of sensitivity can go deeper and deeper for more secured content in the file system. Level 0 (depicted by s0) is the lowest sensitivity level, comparable to say, "public." There can be other sensitivity levels with higher s values: for example, internal, confidential, or regulatory can be depicted by s1, s2, and s3 respectively. This mapping is not stipulated by the policy: system administrators can configure what each sensitivity level mean.

When a SELinux enabled system uses MLS for its policy type (configured in the "/etc/selinux/config" file), it can mark certain files and processes with certain levels of sensitivity. The lowest level is called "current sensitivity" and the highest level is called "clearance sensitivity".

Going hand-in-hand with sensitivity is the category of the resource, depicted by c. Categories can be considered as labels assigned to a resource. Examples of categories can be department names, customer names, projects etc. The purpose of categorization is to further fine-tune access control. For example, you can mark certain files with confidential sensitivity for users from two different internal departments.

For SELinux security contexts, sensitivity and category work together when a category is implemented. When using a range of sensitivity levels, the format is to show sensitivity levels separated by a hyphen (for example, s0-s2). When using a category, a range is shown with a dot in between. Sensitivity and category values are separated by a colon (:). Here is an example of sensitivity / category pair: "user_u:object_r:etc_t:s0:c0.c2". There is only one sensitivity level here and that's s0. The category level could also be written as c0-c2.

So where do you assign your category levels? Let's find the details from the /etc/selinux/targeted/setrans.conf file We won't go into the details of sensitivities and categories here. Just know that a process is allowed read access to a resource only when its sensitivity and category level is higher than that of the resource (i.e. the process domain dominates the resource type). The process can write to the resource when its sensitivity/category level is less than that of the resource.